

Lab 2: Implementing a Shell

COMPSCI 310: Introduction to Operating Systems

1 Shells are cool

Unix [2] embraces the philosophy:

*Write programs that do one thing and do it well.
Write programs to work together.*

The idea of a *shell* command interpreter is central to the Unix programming environment. In particular, the shell executes commands that you enter in response to its prompt in a terminal window, which is called the *controlling terminal*.

```
$ echo Hello
Hello
```

Each command is a character string naming a program or a built-in function to run, followed by zero or more arguments separated by spaces, and (optionally) a few special control directives. If the command names a program to run, then the shell spawns a new child process and executes that program, passing in the arguments. The shell can also direct a process to take its input from a file and/or write its output to a file. It also uses *pipes* to direct the output of one process to the input of another. Thus you can use a shell as an interactive scripting language that combines subprograms to perform more complex functions.

```
$ echo Hello
Hello
# (NOTE: This is a shell comment, just like comments in Java, C, Python, ...)
# echo is a program that simply ‘echoes’ on the terminal whatever it
#   receives as input.
$ echo Hello | wc
1      1      6
# wc is a command that outputs the number of words, lines and bytes in its
#   input; try ‘man wc’ to learn more.
```

Shell can also read and execute a list of commands from a file, called a *shell script*. A shell script is just another kind of program, written in the shell command language—a programming language that is interpreted by the shell. Modern shells support programming language features including variables and control structures like looping and conditional execution. Thus shell scripting is akin to programming in interpreted languages such as Perl, Python, and Ruby.

```
# dumping the contents of a simple shell script
$ cat loop.sh
for i in `seq 1 5`; do
    echo " - $i - "
done
# cat dumps of the contents of its input on the terminal.

# calling the shell ‘sh’ and using it to run the shell script
```

```
$ sh /tmp/loop5.sh
- 1 -
- 2 -
- 3 -
- 4 -
- 5 -
```

There are different kinds of shells—after all they are just programs which can be easily extended to support different features. You can find out which shell you are running by inspecting the `$SHELL` environment variable as follows.

```
$ echo $SHELL
/bin/zsh
```

The shell maintains a set of environment variables with various names and values. This set is essentially a property list of user settings. Each environment variable has a name and a value: both the name and the value are character strings. Shell commands may reference environment variables by name. The environment variable are also passed to all programs that execute as children of the shell.

2 The Devil Shell: dsh

For this lab you will use Unix system calls to implement a simple shell— *devil shell* (*dsh*). `dsh` supports basic shell features: it spawns child processes, directs its children to execute external programs named in commands, passes arguments into programs, redirects standard input/output for child processes, chains processes using pipes, and (optionally) monitors the progress of its children.

You should understand the concepts of environment variables, system calls, standard input and output, I/O redirection, parent and child processes, current directory, pipes, jobs, foreground and background, signals, and end-of-file. These concepts are discussed in class and in the reading.

`dsh` prints a prompt of your choosing (get creative!) before reading each input line. As part of this lab, you will change the prompt to include the process identifier.

```
$ dsh
# starting devil shell from a terminal...

dsh-2400$
# devil shell started; waiting for inputs at the prompt.
# Observe that the PID of dsh - 2400 is part of the prompt
```

`dsh` reads command lines from its standard input and interprets one line at a time. We provide a simple command line parser (in `parse.c`) to save you the work of writing one yourself. You should not need to look at the parser code (really, don't), but you will need to familiarize yourself with the data structures that the parser returns, and how to read important information from them. Your `dsh` must free these structures when it is done with them.

`dsh` exits when the user issues the built-in command `quit`, which you will implement. You can also exit `dsh` by pressing “ctrl-d” at the terminal. “ctrl-d” sends an *end-of-file* marker to signal the program that there is no more input to process (this condition tells `dsh` to quit). The parser already detects the EOF marker and indicates if it was received as input; you merely have to handle this case and route to the control to your `quit` implementation.

```
$ dsh
# starting devil shell from a terminal...

dsh-2353$
# Observe that the pid at the prompt keeps changing.
```

```
dsh-2353$ quit
$
# dsh terminated and control reverts back to the original terminal.
```

The command-line input to **dsh** can contain one or more commands, separated by the special characters - “;” and “|”. The supplied command-line parser supports four special characters “;”, “|”, “<” and “>”. The special characters “<” and “>” are I/O redirection directives: “<” redirects the standard input of a child process to read from a named file, and “>” redirects the standard output of a child process to write to a named file.

We define the syntax of input that **dsh** accepts using a Backus Normal Form (BNF)¹ notation, as follows, to make the shell syntax easy to understand.

Listing 1: BNF for Shell jobs

```

<command-line> ::= <job> (<semicolon> <job>)* [<semicolon>]
<job> ::= <command> (<pipe> <command>)*
<command> ::= <cmd-name> [<arguments>] [<shell-directives>]
<shell-directives> ::= [<less-than> <file-name>] [<greater-than> <file-name>]
<arguments> ::= <argument> (<space> <argument>)* <argument> ::= (<alpha-numeric>)+
```

The semicolon (;) is used to separate the commands on a command line, if there is more than one command. The commands are executed from left to right, sequentially. Unlike regular shells that halt on an error, **dsh** is so cool that it simply ignores the error or failure of a command and moves on to execute the next. The fact that the command failed is recorded in a log file. As part of the lab, you will handle the case where a user input can contain a sequence of commands each terminated by a semicolon (or end-of-line), and execute the commands sequentially in the order specified.

```
dsh-10437$ echo Hello; echo World
10445(Launched): echo Hello
Hello
10446(Launched): echo World
World
```

The special character “|” indicates a pipeline: if a command is followed by a “|” then the shell arranges for its standard output (stdout) to pass through a pipe to the standard input (stdin) of its successor. The command and its successor are grouped in the same *job*. If there is no successor then the command is malformed. Your **dsh** should enable a user to use pipes to compose commands in this way to perform more complex tasks.

```
dsh-22698$ echo Hello | wc -c
22699(Launched): echo Hello | wc -c
6
```

In addition, the last non-blank character of a job may be an “&”. The meaning of “&” is simply that any child processes for the job execute in the background. If a job has no “&” then it runs in the foreground. Implementation of background jobs in **dsh** is **optional**.

Each command (of a job) is a sequence of words (tokens) separated by blank space. The first token names the command: it is either the name of a built-in command or the name of an external program (an executable file) to run. The built-ins (discussed later) are implemented directly within **dsh**.

The **dsh** supports *input/output redirection* using the special characters “<” and “>”. Note that these two special characters are treated as directives to the shell to modify the standard input (**stdin**) and/or standard output (**stdout**) of the command. The directives always follow a command whose standard input and/or output are to be redirected to a file. These concepts are discussed in class and in the reading.

¹BNF: http://en.wikipedia.org/wiki/BackusNaur_Form

- When the special character “<” is used along with a file name, **dsh** (like any regular shell) arranges for the command to read its **stdin** from the specified file.
- When the special character “>” is used along with a file name, **dsh** arranges for the command to write its **stdout** to the specified file.

```
dsh-22698$ echo Hello > hello.txt
22838(Launched): echo Hello > hello.txt
dsh-22698$ wc < hello.txt
22841(Launched): wc < hello.txt
1 1 6
```

2.1 dsh built-ins

The shell executes built-in commands in its own process, without spawning a child process. Whenever an input is supplied, the shell checks for its list of built-in commands—and if the command does not match, the shell spawns a child process for the command to execute. For example, the command **quit** terminates a shell. **quit** is a built-in command which will invoke the **exit()** system call to terminate the shell.

dsh has the following additional built-in commands, which you will implement:

- **jobs**. Output the command strings and status for all active jobs in a numbered sequence. A job is active if it has been launched but has not yet completed, or if it has completed but the shell has not yet reported its completion. Each job has a *job number*, which does not change during the job’s lifetime. A job number is the same as process group number, which is typically the process id of its leader (the first process in the job).
- **fg**. Continue a named job in the foreground. The job is given by a single argument: the job number. If no arguments are specified it continues the last job stopped (or suspended). To stop/suspend a job, a user can press “ctrl+z” in the terminal; you should ensure that suspension works properly in **dsh**.
- **bg**. This is similar to **fg**, but it continues the named job in the background. Implementation of support for **bg** is **optional**.
- **cd**. Change the shell’s current directory. The target is given by a single argument: a pathname.

2.2 Launching external programs

If a command is not recognized as a built-in, then it’s assumed to be an external program to execute. If the command is not an external program either (i.e., the file does not exist), then the error should be logged (discussed in the next section). The program executes in a child process (to **dsh**), grouped with any other child processes that are part of the same job. To implement this part, you will understand and use a few basic Unix system calls (discussed later in section 3) to launch and manage child processes.

2.3 Additional dsh features

Batch mode: The shell should support batch mode (scripting). In batch mode, the shell reads its commands from a file instead of from the controlling terminal. We test your shells in batch mode.

In interactive mode, you display a prompt and the user of the shell will type in one or more commands at the prompt. In batch mode, you specify a batch file on the command line when you start the shell; the batch file contains a sequence of commands as you would have typed them in interactive mode.

In batch mode, you should not display a prompt. You should print each line you read from the batch file back to the user before executing it; this will help you when you debug your shells (and us when we test your programs). When you print the command line, use `fflush` in your C code to ensure that the print completes immediately.

Error handling: The shell prints an error message beginning with “**Error:**” to its `stderr` if an input line is malformed, or on any error condition. Your shell may generate other outputs as well, as you choose.

You can use the `perror()` library function to print informative messages for errors returned by system calls. `perror()` writes a string on the standard error output: the prefix is a string that you specify (“**Error:** ...”), followed by a standard canned summary of the last error encountered during a previous system call or library function. Check also the man pages of `errno`.

When the shell detects that a child process has exited, it should output a string reporting the event and the exit status of the child. In general, the processes of a job will exit together. When the shell detects that a job has exited, the shell frees its job number and discards any record of the job.

3 Notes on System Calls

This section provides some notes on system calls you will need for this lab. It is not a comprehensive guide to these system calls: it is just intended to help by summarizing the most important points. Use the `man` command (or google) to find documentation on the system calls. Note that system call documentation is in “section 2” of the manual; in some cases you must tell `man` which entry you want or it will give you something different. You can tell it by saying, for example, `man 2 pipe`. See also the reading, class notes, and C sample programs.

3.1 Fork

We provide a procedure called `spawn_job` that uses the `fork` system call to create a child process to start a job. The `spawn_job` routine also attends to a few other details. Each job has exactly one *process group* containing all processes in the job, and each process group has exactly one process that is designated as its *leader*. If a job has only one process then that process is the leader of the job’s process group.

Process groups are useful for signal delivery to all the processes in the group. The kernel distributes signals generated from the terminal (ctrl-c, ctrl-z) to all processes in a *foreground* process group. The purpose of the process group is to allow the shell to control which processes receive such signals. Note that `dash` need not and should not declare signal handlers for `ctrl-c` or `ctrl-z`. The lack of support for signals means that `dash` will not notice if one of its background jobs (optional) stops or exits. Instead, the shell has a built-in command called `jobs` to check and report the status of each of its children.

At any given time, at most one process group controls access to the terminal, *tty*, for its input. This process group is the *foreground process group*. The shell sets the foreground process group when starting a new job; by convention the new process group number is the same as the process ID of one of the members of the group, i.e., one of the processes in the job.

A job has multiple processes if it is a pipeline, i.e., a sequence of commands linked by pipes. In that case the first process in the sequence is the leader of the process group. If a job runs in the foreground, then it is the foreground process group and it is bound to the controlling terminal: input typed on the terminal is received by the foreground process group.

The `spawn_job` routine provided with the code shows how to use `fork` and other Unix system calls to set the process group of a process and bind a process group to the controlling terminal using `tcsetpgrp`.

3.2 Exec*

The `exec_()` family of system calls (e.g., `execvp`) enables the calling process to execute an external program that resides in an executable file. An `exec_()` system call **never returns**. Instead, it transfers control to the main procedure of the named program, running within the calling process. All data structures and other state relating to the previous program running in the process—the calling program—are destroyed.

3.3 Wait*

The `wait_()` family of system calls (e.g., `waitpid`) allows a parent process to query the status of a child process or wait for a child process to change state. You use it to implement the `jobs` built-in command, and also to wait for foreground processes to exit or pause (stop), or change state in some other way reported by `waitpid`. The `WNOHANG` option turns `waitpid` into a query: it returns the current status of a child process without actually waiting for child status to change. The `WUNTRACED` option waits until the child exits or stops.

3.4 I/O redirection

Instead of reading and writing from `stdin` and `stdout`, one can choose to read and write from a file. The shell supports I/O redirection using the special characters “<” for input and “>” for output respectively. Redirection is relatively easy to implement: just use `close()` on `stdout` and then `open()` on a file. With file descriptor, you can perform read and write to a file using `creat()`, `open()`, `read()`, and `write()` system calls.

To implement redirection, you should use the `dup2()` system call to duplicate an open file descriptor onto some other specified file descriptor. (See the class notes and reading.)

3.5 Pipelines

A pipeline is a sequence of processes chained by their standard streams, so that the output of each process (`stdout`) feeds directly as input (`stdin`) to the next one. If a command line contains a symbol `|`, the processes are chained by a pipeline.

Pipes can be implemented using the `pipe()` and `dup2()` system calls. A more generic pipeline can be of the form:

```
p1 < inFile | p2 | p3 | .... | pn > outFile
```

where `inFile` and `outFile` are input and output files for redirection.

The descriptors in the child are often duplicated onto standard input or output. The child can then `exec()` another program, which inherits the standard streams. `dup2()` is useful to duplicate the child descriptors to `stdin/stdout`. For example, consider:

```
int fd[2];
pid_t pid;
pipe(fd);

switch (pid = fork()) {
case 0: /* child */
```

```

dup2(fd[0], STDIN_FILENO); /* close stdin (0); duplicate input end of the pipe
                           to stdin */
execve(...);
....
}

```

where `dup2()` closes `stdin` and duplicates the input end of the pipe to `stdin`. The call to `exec_()` will overlay the child's text segment (code) with new executable and inherits standard streams from its parent—which actually inherits the input end of the pipe as its standard input! Now, anything that the original parent process sends to the pipe, goes into the newly exec'ed child process.

3.6 chdir

`chdir` is a simple system call to change the current directory of the calling process to a named directory. See `man 2 chdir`.

4 Getting started

The initial source code is available on the course web. Save the zip file into a directory, unpack it, `cd` into that directory, and type “make”.

4.1 Suggested plan for implementation

1. Read this handout. Read the material from OSTEP on process creation and execution. Bryant and O'Hallaron can be a handy reference [1]. Also see the FAQ on the course web.
2. Read the man pages for `fork()`, `execvp_()`, `wait_()`, `dup2()`, `open()`, `read()`, `write()`, `chdir()`, and `exit()`.
3. Write small programs to experiment with these system calls. There are some programs in the C samples on the course web.
4. Read man pages for `tcsetpgrp()` and `setpgid()`. You should not have to mess with these, but you should know what they do. Read the code we provided for `tcsetpgrp()` and `setpgid()`.
5. Using the parser we gave, start writing single commands. To do this you will need to understand the data structures returned by the parser, but you should not need to look at the parser itself. The parser has a couple of minor peculiarities, but we can live with them. (See the FAQ.) Please do not modify the parser.
6. Add input and output redirection.
7. Add code to display the status of any suspended or background jobs for the `jobs` builtin command. If a background job has completed but has not yet been reported as complete, then `jobs` should report it once. (Support for background jobs is optional.)
8. Add code for error handling and reporting. Check the error returns from the system calls you issue! This will save debugging time and general chaos.
9. Add support for pipeline jobs with two processes using the `pipe()` and `dup2()` system call. Test it.
10. Extend your pipeline support to jobs with three or more processes. Take care. Many groups find this to be strangely confusing.

11. Add the `fg` command.
12. Finish up all of the details.
13. (Optional) Add support for running programs in the background, including `&`, the `bg` command and status notifications for background jobs with the `jobs` command. A “real” shell notifies the user immediately when a background job terminates (how does it know?), but for this lab it is sufficient to report the completion status once in response to the next `jobs` command.
14. Test, test, and test. Be sure your shell works with scripts (batch mode) as well as terminal input!
15. Go *devils*{hell}. Celebrate!

5 What to submit

Submit a single source file named as `dsh.c` along with a README file describing your implementation details, the results, the amount of time you spent on the lab, and also include your name along with the collaborators involved. You can submit the files multiple times before the deadline and we will treat the latest submission as the final submission for grading.

The grading is done on the scale of 100 as per below:

- Basic command support (process creation and execution) : 20 points
- Input/Output redirection and built-in command `cd` : 15 points
- Two-process pipeline jobs: 20 points
- Pipeline jobs of more than two processes: 15 points
- Process groups, job control (`ctrl-z` and `fg`), `jobs` command, error handling, pid prompt, other: 15 points
- README : 15 points
- Background jobs: 20 points extra credit

References

- [1] Randal E. Bryant and David R. O'Hallaron. *Exceptional Control Flow, Excerpts from Chapter 8 from Computer Systems: A Programmer's Perspective, 2/E (CS:APP2e)*. <http://www.cs.duke.edu/courses/fall12/compsci210/internal/controlflow.pdf>.
- [2] Dennis M. Ritchie and Ken Thompson. The unix time-sharing system. *Commun. ACM*, 17(7):365–375, July 1974.

A dsh by example

Here we provide a walk-through of the *devil shell* by issuing some sample jobs. You can issue all the commands yourself on a linux system using the shell, `dsh-example`, available in the repository. (Or use `dsh-osx` on Mac systems.)

Note: the example `dsh` logs all errors to a file instead of printing them. This used to be a requirement of the lab, but it was removed as a requirement. You may see stray references to it here and there. This behavior is acceptable but not required.

Note: The extra line after each job execution is added here for the readability and need not be present in the actual `dsh` output.

```
# starting dsh
$ ./dsh-example

# Observe the pid of dsh in the prompt.
dsh-23822$

# listing current directory contents
dsh-23822$ ls
23823(Launched): ls
dsh.c  dsh-example  dsh.h  dsh.log  Makefile

# The above job consists of one command 'ls' which is an external program.
# \dsh spawns a child process and launches this external program in the child.
# The programs runs, prints its output and quits.
# This finished job can be viewed using the jobs (built-in) command.

# checking the status of the finished 'ls' job
dsh-23822$ jobs
23823(Completed): ls

# completed jobs are only displayed once.

dsh-23822$ jobs

# now, jobs shows nothing!

# launching multiple jobs using a semicolon
dsh-23822$ ls; ps
23840(Launched): ls
dsh.c  dsh-example  dsh.h  dsh.log  Makefile
23841(Launched): ps
PID TTY          TIME CMD
23841 pts/13        00:00:00 ps
17690 pts/13        00:00:00 tcsh
17705 pts/13        00:00:00 bash
23822 pts/13        00:00:00 dsh-example

# note that 'ps' ran after 'ls'

# Note that semicolon is used as a delimiter to specify multiple jobs
dsh-23822$ jobs
23840 (Completed): ls
23841(Completed): ps

# Observe that the output above shows two jobs (issued in the same line)

# pipes
dsh-23822$ ls | wc -l
23843(Launched): ls | wc -l
5
```